

Chef was invited to the party of  $N$  people numbered from  $1$  to  $N$ . Chef knows the growth of all the people, i.e. he knows the growth of the  $i^{\text{th}}$  person is denoted by an integer  $A_i$  not exceeding  $M$ .

Chef decided to have some fun. At first, he forms  $K$  groups of people. The  $i^{\text{th}}$  group consists of all the people numbered from  $L_i$  to  $R_i$ . Groups may overlap too.

For each group, Chef wants to know the following information: the maximum difference between the numberings of two people having same growth. Formally, Chef wants to know the following:

$$\max\{|x - y| : L_i \leq x, y \leq R_i \text{ and } A_x = A_y\}$$

Please help Chef to have fun.

## Input

There is only one test case in one test file.

The first line of input contains three space-separated integers  $N$ ,  $M$  and  $K$ , denoting the number of people, the maximum growth and number of groups respectively. The second line contains  $N$  space-separated integers  $A_1, A_2, \dots, A_N$  denoting the growth of people. Then the  $i^{\text{th}}$  line of the next  $K$  lines contains two space-separated integers  $L_i, R_i$ , denoting the  $i^{\text{th}}$  group.

## Output

For each group, output the integer denoting the maximum difference between numbering of two people having same growth in a single line.

## Constraints and Subtasks

- $1 \leq A_i \leq M$
- $1 \leq L_i \leq R_i \leq N$

### Subtask 1: 20 points

- $1 \leq N, M, K \leq 1000 = 10^3$

### Subtask 2: 80 points

- $1 \leq N, M, K \leq 100000 = 10^5$

## Example

### Input:

```
7 7 5
4 5 6 6 5 7 4
```

6 6  
5 6  
3 5  
3 7  
1 7

**Output:**

0  
0  
1  
1  
6

**Explanation**

**Group 1.** There is only one person in the group. Thus the maximum difference of numbers should be **0**.

**Group 2.** There are two persons in the group. Their growth are  $A_5 = 5$  and  $A_6 = 7$ . Thus there is no pair of persons who have the same growth. Thus the answer for this group will also be **0**.

**Group 3.** There are three persons in the group. Their growth are  $A_3 = 6$ ,  $A_4 = 6$  and  $A_5 = 5$ . Here person **3** and person **4** has the same growth. Thus the answer is  $|4 - 3| = 1$ .

**Group 4.** There are more persons than the group **3**. But they has different growth, other than person **3** and person **4**. Thus the answer is also  $|4 - 3| = 1$ .

**Group 5.** This group contains all the people, and person **1**and person **7** has the same growth  $A_1 = A_7 = 4$ . So the answer is  $|7 - 1| = 6$ .

**DIFFICULTY:**

MEDIUM

**PREREQUISITES:**

sqrt decomposition, preprocessing

**PROBLEM:**

Given a sequence of  $N$  integers  $A_1, A_2, \dots, A_N$ , where each  $A_i$  is between 1 to  $M$ , you are to answer  $Q$  queries of the following kind:

- Given  $L$  and  $R$ , where  $1 \leq L \leq R \leq N$ , what is the maximum  $|x - y|$  such that  $L \leq x, y \leq R$  and  $A_x = A_y$ ?

Note that in the problem,  $Q$  is actually  $K$ .

**QUICK EXPLANATION:**

For each  $i$ ,  $1 \leq i \leq N$ , precompute the following in  $O(N)$  time:

- $\text{next}[i]$ , the smallest  $j > i$  such that  $A_i = A_j$
- $\text{prev}[i]$ , the largest  $j < i$  such that  $A_i = A_j$

Let  $S = \lfloor \sqrt{N} \rfloor$ , and  $B = \lceil N/S \rceil$ . Decompose the array into  $B$  blocks, each of size  $S$  (except possibly the last). For each  $i$ ,  $1 \leq i \leq N$ , and  $0 \leq j \leq B - 1$ , precompute the following in  $O(N\sqrt{N})$  time:

- $\text{last\_in\_blocks}[j][i]$ , the largest  $k \leq jS + S$  such that  $A_k = A_i$
- $\text{block\_ans}[j][i]$ , the answer for the query  $(L, R) = (jS + 1, i)$ . For a fixed  $j$ , all the  $\text{block\_ans}[j][i]$  can be computed in  $O(N)$  time.

Now, to answer a query  $(L, R)$ , first find the blocks  $j_L$  and  $j_R$  where  $L$  and  $R$  belong in  $(0 \leq j_L, j_R < B)$ . Then the answer is at least  $\text{block\_ans}[j_L + 1][R]$ , and the only pairs  $(x, y)$  not yet considered are those where  $L \leq x \leq j_L S + S$ . To consider those, one can simply try all  $x$  in that range, and find the highest  $y \leq R$  such that  $A_x = A_y$ . Finding that  $y$  can be done by using  $\text{last\_in\_blocks}[j_R - 1][x]$  and a series of  $\text{next}$  calls. To make that last part run in  $O(S)$  time, consider only the  $x$  such that  $\text{prev}[x] < L$ .

**EXPLANATION:**

We'll explain the solution for subtask 1 first, because our solution for subtask 2 will build upon it. However, we will first make the assumption that  $M \leq N$ , otherwise we can simply replace the values  $A_1, \dots, A_N$  with numbers from 1 to  $N$ , and should only take  $O(N)$  time with a set. However, we don't recommended that you actually do it; this is only to make the analysis clearer.

 **$O(N^2)$  per query**

First, a simple brute-force  $O(N^2)$ -time per query is very simple to implement, so getting the first subtask is not an issue at all. I'm even providing you with a pseudocode on how to do it :)

```
def answer_query(L, R):
    for d in R-L...1 by -1
        for x in L...R-d
            y = x+d
            if A[x] == A[y]
                return d

    return 0
```

We're simply checking every possible answer from  $[0, R - L]$  in decreasing order. Note that the whole algorithm runs in  $O(QN^2)$  time, which *could* get TLE if the test cases were stronger. But in case you can't get your solution accepted, then it's time to optimize your query time to...

## $O(N)$ per query

To obtain a faster running time, we have to use the fact that we are finding the maximum  $|x - y|$ . What this means is that for every value  $v$ , we are only concerned with the first and last time it occurs in  $[L, R]$ .

We first consider the following alternative  $O(N^2)$ -time per query solution:

```
def answer_query(L, R):
    answer = 0
    for y in L...R
        for x in L...y
            if A[x] == A[y]
                answer = max(answer, y - x)

    return answer
```

The idea here is that for every  $y$ , we are seeking  $A_x$ , which is the first occurrence of  $A_y$  in  $[L, y]$ , because all the other occurrences will result in a smaller  $y - x$  value. Now, to speed it up, notice that we don't have to recompute this  $x$  every time we encounter the value  $A_x$ , because we are already reading the values  $A_L, \dots, A_R$  in order, so we already have the information "when did  $A_y$  first appear" before we ever need it! Here's an implementation (in pseudocode):

```
def answer_query(L, R):
    index = new map/dictionary
    answer = 0
    for y in L...R
        if not index.has_key(A[y])
            index[A[y]] = y
        answer = max(answer, y - index[A[y]])

    return answer
```

Now, notice that this runs in  $O(N)$  time if one uses a hash map for example!

We mention here that it's possible to drop the use of a hash map by using the fact that the values  $A_y$  are in  $[1, M]$ . This means that we can simply allocate an *array* of length  $M$ , instead of creating a hash map from scratch or clearing it. However, we must be careful when we reinitialize this array, because it is long! There are two ways of "initializing" it:

- We clear the array every time we're done using it, but we only clear those we just encountered. This required listing all the indices we accessed.
- We maintain a *parallel array* that contains when array was last accessed for each index. To clear the array, we simply update the *current time*.

We'll show how to do the second one:

```
class LazyMap:
    index[1..M]
    found[1..M] # all initialized to zero
    time = 0

    def clear():
        this.time++

    def has_key(i):
        return this.found[i] == this.time

    def set(i, value): # called on the statement x[i] = value for example
        this.found[i] = this.time
```

```

this.index[i] = value

def get(i): # called on the expression x[i] for example
    return this.index[i]

index = new LazyMap()

def answer_query(L, R):
    index.clear()
    answer = 0
    for y in L...R
        if not index.has_key(A[y])
            index[A[y]] = y
        answer = max(answer, y - index[A[y]])

    return answer

```

Using this, the algorithm still runs in  $O(N)$  time (remember that we assume  $M \leq N$ ), but most likely with a lower constant.

The overall algorithm runs in  $O(QN)$  time.

## sqrt decomposition

When one encounters an array with queries in it, there are usually two ways to preprocess the array so that the queries can be done in sublinear time:

- **sqrt decomposition**, which splits up the array into  $\lceil N/S \rceil$  blocks of size  $S$  each.  $S$  is usually taken to be  $\lfloor \sqrt{N} \rfloor$  (hence the term "sqrt decomposition"). Usually, one can reduce the running time to  $O((N + Q)\sqrt{N})$  or  $O((N + Q)\sqrt{N \log N})$ . Sometimes, depending on the problem, it may also yield  $O((N + Q)N^{2/3})$  time.
- **build some tree structure on top of the array**. This usually yields an  $O(N + Q \log N)$  or  $O((N + Q) \log N)$  time algorithm.

There are other less common ways, such as **lazy updates** or combinations of the above, but first we'll try out whether the above work.

Suppose we have selected the parameter  $S$ , and we have split the array into  $B = \lceil N/S \rceil$  blocks of size  $S$ , except possibly the last block which may contain fewer than  $S$  elements. Suppose we want to answer a particular query  $(L, R)$ . Note that  $L$  and  $R$  will belong to some block. For simplicity, we assume that they belong to different blocks, because if they are on the same block, then  $R - L \leq S$ , so we can use the  $O(S)$  time query above.

Thus, the general picture will be:



We have marked two additional points,  $E_L$  and  $E_R$ , which are the boundaries of the blocks completely inside  $[L, R]$ . Now, it would be nice if we have already precomputed the answer for the query pair  $(E_L, E_R)$ , because then we will only have to deal with at most  $2(S - 1)$  remaining values:  $[L, E_L]$  and  $[E_R, R]$ . We can indeed precompute the answers at the boundaries, but we can do even better: we can precompute the answers for all pairs  $(E, R)$ , where  $E$  is a boundary point and  $R$  is *any* point in the array! There are only  $O(BN)$  pairs, and we can compute the answers in  $O(BN)$  time also:

```

class LazyMap:
    ...

S = floor(sqrt(N))
B = ceil(N/S)
index = new LazyMap()

```

```

block_ans[1..B][1..N]
def precompute():
    answer = 0
    for b in 1..B
        index.clear()
        E = b*S-S+1 # left endpoint of the b'th block
        answer = 0
        for R in E..N
            if not index.has_key(A[R])
                index[A[R]] = R
            answer = max(answer, R - index[A[R]])
        block_ans[b][R] = answer

```

(if you read the "quick explanation", note that there is a slight difference here: we're indexing the blocks from 1 to B instead of 0 to  $B - 1$ )

This means that, in the query, the only remaining values we haven't considered yet are those in  $[L, E_L]$ . To consider those, we have to know, for each  $x$  in  $[L, E_L]$ , the last occurrence of  $A_x$  in  $[L, R]$ . To do so, we will need the following information:

- $\text{next}[i]$ , the smallest  $j > i$  such that  $A_i = A_j$
- $\text{prev}[i]$ , the largest  $j < i$  such that  $A_i = A_j$
- $\text{last\_in\_blocks}[j][i]$ , the largest  $k$  within the first  $j$  blocks such that  $A_k = A_i$

How will this help us? Well, we want to find  $A_x$ 's last occurrences in  $[L, R]$ . So first, we find its last occurrence in the blocks up to  $E_R$  (it's just  $\text{last\_in\_blocks}[\text{floor}(R/S)][x]$ ). However, it's possible that  $A_x$  appears in  $[E_R, R]$ , so we need to use its next pointers, until we find the *last* one. Since there are at most  $S - 1$  elements in  $[E_R, R]$ , this *seems* fast, but it could easily take  $O(S^2)$  time for example when most of the values in  $[L, E_L]$  and  $[E_R, R]$  are equal. Thankfully, this is easily fixed: we only care about the *first* occurrence of  $A_x$ , so if it has been encountered before, then we don't have to process it again! This ensures that for *distinct* value in  $[E_R, R]$ , its set of indices is iterated only once. This therefore guarantees an  $O(S)$  running time!

Checking whether an  $A_x$  has been encountered before can also be done using the `index` approach, or alternatively as  $\text{prev}[x] \geq L$ :

```

def answer_query(L, R):
    b_L = ((L+S-1)/S)
    b_R = R/S
    if b_L >= b_R
        # old query here
    else
        E_L = b_L*S
        answer = block_ans[b_L+1][R]
        for x in L..E_L
            if prev[x] < L # i.e. x hasn't been encountered before
                y = last_in_blocks[floor(R/S)][x]
                while next[y] <= R
                    y = next[y]
                answer = max(answer, y - x)

    return answer

```

One can now see that the query time is now  $O(S)$  :) Note that  $b_L \leq b_R$  means that  $L$  and  $R$  are within  $O(S)$  elements away, so we can do the old query instead.

Let's now see how to precompute `next`, `prev` and `last_in_blocks`. First, `next[i]` and `prev[i]` can easily be computed in  $O(N)$  time with the following code:

```

...
next[1..N]
prev[1..N]
last[1..M] # initialized to 0

```

```

...
def precompute():
    ...

    for i in 1..N
        next[i] = N+1
        prev[i] = 0

    for i in 1..N
        j = last[A[i]]
        if j != 0
            next[j] = i
            prev[i] = j
        last[A[i]] = i

```

The `last` array stores the last index encountered for every value, and is updated as we traverse the array.

And then `last_in_blocks` can be compute in  $O(BN)$  time:

```

...
last_in_blocks[1..B][1..N] # initialized to 0
...

def precompute():
    ...
    for b in 1..B
        L = b*S-S+1
        R = min(b*S,N)
        for y in L..R
            if next[y] > R
                x = y
                while x > 0
                    last_in_blocks[b][x] = y
                    x = prev[x]

        for x in 1..N
            for b in 2..B
                if last_in_blocks[b][x] == 0
                    last_in_blocks[b][x] = last_in_blocks[b-1][x]

```

The first loop finds the last value encountered at each block (with the check  $next[y] > R$ ), and proceeds setting the `last_in_blocks` of all the indices until that position with equal value, using the `prev` pointer. The second loop fills out the remaining entries, because some values do not have representatives in some blocks.

## Running time

Now, what is the total running time then? The precomputation runs in  $O(BN)$  time, and each query takes  $O(S)$  time, so overall it is  $O(NB + QS)$ . But remember that  $B = \Theta(N/S)$ , so the algorithm is just  $O(N^2/S + QS)$ . But we still have the freedom to choose the value of  $S$ . Now, most will simply choose  $S = \Theta(\sqrt{N})$ , so that the running time is  $O((N + Q)\sqrt{N})$ , but we are special, so we will be more pedantic.

Note that  $N^2/S$  is a decreasing function while  $QS$  is an increasing function. Also, remember that  $O(f(x) + g(x)) = O(\max(f(x), g(x)))$  (why?). Therefore, the best choice for  $S$  is one that makes  $N^2/S$  and  $QS$  equal (at least asymptotically). Thus, we want the choice  $S = \Theta(N/\sqrt{Q})$  instead, and the running time is  $O(N\sqrt{Q} + Q)$  (the  $+Q$  is there to account for when  $Q > N^2$ ). For this problem, there's not much difference between this and  $O((N + Q)\sqrt{N})$ , but the running time  $O(N\sqrt{Q} + Q)$  is mostly for theoretical interest, and when  $Q$  is much less than  $N$  (or much more), you'll feel the difference.

Optimization side note: there is another way to do the old query without using our `LazyMap`, or at least calling `has_key`: *traverse the array backwards*. Here is an example:

```

...
_index[1..M]
...

def answer_query(L, R):
    ...
    if b_L >= b_R
        # old query
        answer = 0
        for y in R...L by -1
            _index[A[y]] = y

        for y in L...R
            answer = max(answer, y - _index[A[y]])
    else
        ...

    return answer

```

I found that this is a teeny teeny bit faster than the original  $O(S)$  old query :)

Also, when choosing  $S$ , one does not have to choose  $\lfloor \sqrt{N} \rfloor$ , or even  $\lfloor N/\sqrt{Q} \rfloor$ , because there is still a constant hidden in the  $\Theta$  notation. This means that you still have the freedom to choose a multiplicative constant for  $S$ , which in practice essentially amounts to the freedom to select  $S$  however you want. To get the best value for  $S$ , try generating a large input (with varying values of  $M$  !), and finding the best choice for  $S$  via [ternary search](#). The goal is to get the precomputation part and the query part roughly equal in running time. This technique of **tweaking the parameters** is incredibly useful in long contests where the time limit is usually tight.

## Time Complexity:

$O(M + (N + Q)\sqrt{N})$  but theoretically it is  $O(N\sqrt{Q} + Q)$

Note that in the problem,  $Q$  is actually  $K$ .

```

1 #include<bits/stdc++.h>
2 using namespace std;
3 const int MAXN = 1e5+5;
4 const int MAXV = 1e5+5;
5 const int SQRN = 350;
6 int N, M, Q, BLOCKN, BLOCK_SIZE, a[MAXN];
7 int Prev[MAXN], Next[MAXN], last[MAXV], indexx[MAXV], found[MAXV];
8 int last_in_blocks[SQRN][MAXN], block_ans[SQRN][MAXN];
9
10 void preprocess(){
11
12     BLOCK_SIZE = floor(sqrt(N));
13     BLOCKN = ceil(N/BLOCK_SIZE);
14
15     for(int i=1; i<=M; i++) last[i] = 0;
16     for(int i=1; i<=N; i++) Prev[i] = 0;
17     for(int i=1; i<=N; i++) Next[i] = N+1;
18
19     for(int i=1; i<=N; i++){
20         if(last[a[i]] != 0){
21             Next[last[a[i]]] = i;
22             Prev[i] = last[a[i]];
23         }
24         last[a[i]] = i;
25     }
26
27     memset(last_in_blocks, 0, sizeof(last_in_blocks));
28     for(int b=1; b<=BLOCKN; b++){
29         int L = (b*BLOCK_SIZE) - BLOCK_SIZE + 1;
30         int R = min(b*BLOCK_SIZE, N);
31         for(int y=L; y<=R; y++){
32             if(Next[y]>R){
33                 int x = y;
34                 while(x>0){
35                     last_in_blocks[b][x] = y;
36                     x = Prev[x];
37                 }
38             }
39         }
40     }
41     for(int x=1; x<=N; x++){
42         for(int b=2; b<=BLOCKN; b++){
43             if(last_in_blocks[b][x]==0){
44                 last_in_blocks[b][x] = last_in_blocks[b-1][x];
45             }
46         }
47     }
48
49     int mytime = 0;
50     memset(found, 0, sizeof(found));
51     memset(indexx, 0, sizeof(indexx));
52     for(int b=1; b<=BLOCKN; b++){
53         mytime++;
54         int ans = 0;
55         int L = (b*BLOCK_SIZE) - BLOCK_SIZE + 1;
56         for(int y=L; y<=N; y++){
57             if(found[a[y]]!=mytime){
58                 indexx[a[y]] = y;
59                 found[a[y]] = mytime;
60             }
61             ans = max(ans, y-indexx[a[y]]);
62             block_ans[b][y] = ans;
63         }
64     }
65 }
66 }
```

```
67 int queryInBlock(int L,int R){  
68     map<int,int>MP;  
69     int ans = 0;  
70     for(int x=L; x<=R; x++){  
71         if(MP.find(a[x])==MP.end()){  
72             MP[a[x]] = x;  
73         }  
74         ans = max(ans, x-MP[a[x]]);  
75     }  
76     return ans;  
77 }  
78  
79 int solve(int L, int R){  
80     int Lblock = (L+BLOCK_SIZE-1) / BLOCK_SIZE;  
81     int Rblock = (R+BLOCK_SIZE-1) / BLOCK_SIZE;  
82     int ans = 0;  
83     if(Lblock==Rblock || Lblock+1==Rblock){  
84         ans = queryInBlock(L,R);  
85     }  
86     else{  
87         int LblockEnd = Lblock*BLOCK_SIZE;  
88         ans = block_ans[Lblock+1][R];  
89         for(int x=L; x<=LblockEnd; x++){  
90             if(Prev[x]<L){  
91                 int y = last_in_blocks[Rblock-1][x];  
92                 while(Next[y]<=R){  
93                     y = Next[y];  
94                 }  
95                 ans = max(ans, y-x);  
96             }  
97         }  
98     }  
99     return ans;  
100 }  
101  
102 int main(){  
103     scanf("%d%d%d",&N,&M,&Q);  
104     for(int i=1; i<=N; i++){  
105         scanf("%d",&a[i]);  
106     }  
107     preprocess();  
108  
109     for(int i=1; i<=Q; i++){  
110         int L,R; scanf("%d%d",&L,&R);  
111         if(L>R)swap(L,R);  
112         int ans = solve(L,R);  
113         printf("%d\n",ans);  
114     }  
115     return 0;  
116 }  
117 }  
118 }
```